



SIGGRAPH2015
Xroads of Discovery

The 42nd International Conference and Exhibition
on Computer Graphics and Interactive Techniques

The Rendering Pipeline - Challenges & Next Steps

JOHAN ANDERSSON

ELECTRONIC ARTS

Intro

- ▶ What does an advanced game engine **real-time rendering pipeline** look like?
- ▶ What are some of the key **challenges & open problems**?
- ▶ What are some of the **next steps** to improve on?
- ▶ From both **software & hardware** perspectives

Previous talks



Beyond Programmable Shading Course
ACM SIGGRAPH 2010

5 Major Challenges in Interactive Rendering

Johan Andersson
DICE




Beyond Programmable Shading Course
ACM SIGGRAPH 2012

~~25~~ MAJOR CHALLENGES IN REAL-TIME RENDERING

Johan Andersson, DICE

Beyond Programmable Shading, SIGGRAPH 2012

2010 & 2012 challenges



Long term goal:

Photo-realistic rendering at 1W

Improvements since 2010 & 2012

- ▶ Image quality & authoring: massive transition to PBR
- ▶ Reflections: SSR and perspective-correct IBLs
- ▶ Antialiasing: TAA instead of MSAA
- ▶ Gen4 consoles (PS4 & XB1) as new minspec
- ▶ Compute shader use prevalent – create your own pipelines!

Improvements since 2010 & 2012 (cont.)

- ▶ New **explicit control APIs**
 - ▶ Mantle, Metal, DX12, Vulkan
 - ▶ Well needed change & major step forward
 - ▶ Not much improvements on compute & shaders
- ▶ **Programmability**
 - ▶ Conservative raster, min/max texture filter
 - ▶ "Need a virtual data-parallel ISA" -> **SPIR-V!**
 - ▶ "Render target read/modify/write" -> **Raster Ordered Views**
 - ▶ Sparse resources

Pipeline of today – key themes

- ▶ Non-orthogonality gets in the way – can we get to a more unified pipeline?
- ▶ Complexity is continuing to increase
- ▶ Increasing quality in a scalable way

Getting to a more unified pipeline

Transparencies – sorting

- ▶ Can't mix different transparent surfaces & volumes
 - ▶ Particles, meshes, participating media, raymarching
 - ▶ Can't render strict front to back to get correct sorting
 - ▶ Most particles can be sorted, have to use **uber shaders**
- ▶ Constrains **game environments**
- ▶ Restricts games from using more **volumetric rendering**



particles



meshes



volumetric

Transparencies – sorting solution

- ▶ Render everything with **Order-Independent Transparency** (OIT)
 - ▶ Use Raster Ordered Views (DX12 FL: Haswell & Maxwell)
 - ▶ Not available on consoles = most games stuck with no OIT
- ▶ Scalable to **mix all types** of transparencies with high quality?
 - ▶ Transparent meshes (windows, foliage): 1-50x overdraw
 - ▶ Particles: 10-200x overdraw
 - ▶ Volume rendering (ray-marched)
- ▶ Able to combine with **variable resolution rendering**?
 - ▶ Most particles & participating do not need to be shaded at full resolution

Defocus & motion blur - opaque

- ▶ Works okay in-game on **opaque surfaces**
 - ▶ Render out velocity vectors
 - ▶ Calc CoC from z
 - ▶ Apply post-process
- ▶ But **not correct** or ideal
 - ▶ Leakage
 - ▶ Disocclusion



Defocus & motion blur - transparencies

- ▶ Transparent surfaces are even more **problematic**
 - ▶ Esp. motion blur
 - ▶ Typically simulated with **stretched geometry** (works mostly for sparks)
 - ▶ Can only skip or smear everything with standard **post-processes**
 - ▶ Fast moving particles should also have internal motion blur



Defocus & motion blur - transparencies

- ▶ Blend velocity vectors & CoC for transparencies?
 - ▶ Feed into the post-process passes
 - ▶ Post-processes should also be depth-aware - use OIT approx. transmittance function
 - ▶ Still not correct, but could prevent the biggest artifacts
- ▶ Ideal: directly sample defocus & motion blur in rendering
 - ▶ But how? Stochastic raster? Raytracing?
 - ▶ Pre-filtered volumetric representations?

Forward vs deferred

- ▶ Most high-end games & engines use **deferred shading** for opaque geometry
 - ▶ Better quad utilization
 - ▶ Separation of material property laydown & lighting shaders
- ▶ Would like to **render more as transparent**, which has to use forward
 - ▶ Thin geometry: hair & fur
 - ▶ Proxy geometry (foliage) with alpha-blending for antialiasing
- ▶ But forward rendering is much more **limiting in compositing**
 - ▶ No SSAO
 - ▶ No screen-space reflections
 - ▶ No decal blending of individual channels (e.g. albedo)
 - ▶ No screen-space sub-surface scattering

Forward vs deferred (cont.)

- ▶ Can we extend either forward or deferred to be **more orthogonal**?
- ▶ Use **world-space** data structures instead of **screen-space**
 - ▶ Be able to query & calc AO, reflections, decals while forward shading
 - ▶ **Texture shader** to convolve SSS lighting
 - ▶ **Massive** forward **uber shaders** that can do everything
- ▶ Render opaque & transparent with **deep deferred shading**?
 - ▶ Store all layers of a pixel, including transparents, in a deep gbuffer
 - ▶ Unbounded memory
 - ▶ Be able to query neighbors (AO)
 - ▶ Be able to render into with blending (decals)

Rendering pipeline complexity

Rendering pipeline complexity

- ▶ Recent improvements that **reduce complexity** 😊
 - ▶ **New APIs** are more explicit – less of a black box
 - ▶ **DX11 hardware & compute shaders** now minspec
 - ▶ Hardware trend towards **DX12 feature level**

Rendering pipeline complexity

- ▶ Challenges:
 - ▶ Sheer **amount** of rendering systems & passes
 - ▶ Making **architectural choices** of what techniques & pipeline to use
 - ▶ Shader permutations & **uber shaders**
 - ▶ **Compute shaders** still very limiting – no nested dynamic parallelism & pipes
 - ▶ Mobile: **TBDR** vs immediate mode



Battlefield 4 rendering passes

- ▶ reflectionCapture
- ▶ planarReflections
- ▶ dynamicEnvmap
- ▶ mainZPass
- ▶ **mainGBuffer**
- ▶ mainGBufferSimple
- ▶ mainGBufferDecal
- ▶ decalVolumes
- ▶ mainGBufferFixup
- ▶ msaaZDown
- ▶ msaaClassify
- ▶ lensFlareOcclusionQueries
- ▶ lightPassBegin
- ▶ **cascadedShadowmaps**
- ▶ spotlightShadowmaps
- ▶ downsampleZ
- ▶ linearizeZ
- ▶ ssao
- ▶ hbaoHalfZ
- ▶ hbao
- ▶ ssr
- ▶ halfResZPass
- ▶ **halfResTransp**
- ▶ mainDistort
- ▶ **lightPassEnd**
- ▶ mainOpaque
- ▶ linearizeZ
- ▶ mainOpaqueEmissive
- ▶ mainTransDecal
- ▶ fgOpaqueEmissive
- ▶ subsurfaceScattering
- ▶ skyAndFog
- ▶ hairCoverage
- ▶ mainTransDepth
- ▶ linerarizeZ
- ▶ mainTransparent
- ▶ halfResUpsample
- ▶ motionBlurDerive
- ▶ motionBlurVelocity
- ▶ motionBlurFilter
- ▶ filmicEffectsEdge
- ▶ spriteDof
- ▶ fgTransparent
- ▶ lensScope
- ▶ filmicEffects
- ▶ bloom
- ▶ luminanceAvg
- ▶ **finalPost**
- ▶ overlay
- ▶ fxaa
- ▶ smaa
- ▶ resample
- ▶ screenEffect
- ▶ hmdDistortion

Architectural decisions

- ▶ Selecting **which techniques** to develop & invest in is a challenge
 - ▶ Critical to create visual look of a game
 - ▶ Non-orthogonal choices and tradeoffs
 - ▶ Difficult to predict the moving future of hardware, games and authoring
- ▶ Can be **paralyzing** with a big advanced engine rendering pipeline
 - ▶ Exponential scaling with amount of systems & techniques interacting
 - ▶ Difficult to redesign and move large passes
 - ▶ Can result in a lot of refactoring & cascading effects to the overall pipeline
 - ▶ Backwards compatibility with existing content
- ▶ Easier if passes & systems can be made more **decoupled**

What can we do to reduce complexity?

- ▶ A more **unified pipeline** would certainly help!
 - ▶ Such as with OIT
 - ▶ Or in the long term: native handling of defocus & motion blur
- ▶ Improve **GPU performance** – simplify rendering systems
 - ▶ Much of the complexity comes from optimizations for performance
 - ▶ Could sacrifice a bit of performance for increased orthogonality, but not much
 - ▶ We have real-time constrain = get the most out of our 16 ms/f (VR: 4 ms/f!)
- ▶ **Raytrace & raymarch** more
 - ▶ Easier to express complex rendering
 - ▶ Warning: moves to complexity to data structures and the GPU execution instead
 - ▶ Not practical overall replacement / unification
 - ▶ Use as complement – more & more common (SSR, volume rendering, shadows?)

What can we do to reduce complexity?

- ▶ Make it **easier to drive** the graphics & compute
 - ▶ CPU/GPU communication – C++ on both sides (and more languages)
 - ▶ Device enqueue & nested data parallelism
 - ▶ Increase flexibility, expressiveness & modularity of building pipelines
- ▶ Build a **specialized renderer**
 - ▶ Focus in on very specific rendering techniques & look
 - ▶ Typically tied to a single game
 - ▶ E.g. The Tomorrow Children, Dreams
- ▶ Build engines, tools & infrastructure to build **general renderers**
 - ▶ Handle wide set of environments, content and techniques
 - ▶ Modular layers to easily have all the passes & techniques interoperate
 - ▶ Shader authoring is also key

Uber shaders

- ▶ Example cases:
 - ▶ **Forward shaders** (lights, fog, skinning, etc)
 - ▶ **Particles** (to be able to sort without OIT) – want to use individual shaders instead
 - ▶ **Terrain layers** [Andersson07] – want to use massive uber shaders
- ▶ Why they can be a problem:
 - ▶ **Authoring**: Massive shaders with all possible paths in it, no separate shader linker
 - ▶ **Performance**: Large GPR pressure affects entire shader
 - ▶ **Performance**: Flow control overhead
- ▶ Classic approach: break out into **separate shader permutations**
 - ▶ Static CPU selection of shader/PSO to use – limited flexibility
 - ▶ Can end up creating huge amount of permutations = long compile/load times.
 - ▶ Worse with new APIs! PSO explosion

Uber shaders – potential improvements

- ▶ Shader **function pointers**
 - ▶ Define individual functions as own kernels
 - ▶ Select pointers to use per draw call
 - ▶ Part of ExecuteIndirect params
 - ▶ Ideal: Select pointers inside shader – not possible today
 - ▶ Optimization: VS selects pointers PS will use?
 - ▶ What would the **consequences** be for the GPU?
 - ▶ I\$ stalls, register allocation, coherency, more?
- ▶ More efficient **GPU execution** of uber shaders?
 - ▶ Shaders with highly divergent flow & sections with very different GPR usage
 - ▶ Hardware & execution model that enables resorting & building **coherency**?

Scalable quality




SPEED HUNTERS YUTA NAKAMURA

**NEED FOR
SPEED**

Real-time rendering have gotten quite far!

- ▶ In order to get further, want to:
 1. Get that last 5-10% *quality* in our environments to reach photorealism



NFS photo reference

Real-time rendering have gotten quite far!

► In order to get further, want to:

1. Get that last **5-10% quality** in our environments to reach photorealism
2. Be able to build & render **new environments** that we haven't been able to before



Glass houses!

Real-time rendering have gotten quite far!

- ▶ In order to get further, want to:
 1. Get that last 5-10% **quality** in our environments to reach photorealism
 2. Be able to build & render **new environments** that we haven't been able to before



Dreams (MediaMolecule)

Difficult areas

- ▶ Hair & fur
 - ▶ OIT, overdraw, LOD, quad overshading, deep shadows
- ▶ Foliage
 - ▶ OIT, overdraw, LOD, geometry throughput,
 - ▶ Lighting, translucency, AO
- ▶ Fluids
 - ▶ LOD & scalability, simulation, overall rendering
- ▶ VFX
 - ▶ Need volumetric representation & lighting
 - ▶ Related to [Hillaire15]



Difficult areas

- ▶ **Hair & fur**
 - ▶ OIT, overdraw, LOD, quad overshading, deep shadows
- ▶ **Foliage**
 - ▶ OIT, overdraw, LOD, geometry throughput,
 - ▶ Lighting, translucency, AO
- ▶ **Fluids**
 - ▶ LOD & scalability, simulation, overall rendering
- ▶ **VFX**
 - ▶ Need volumetric representation & lighting
 - ▶ Related to [Hillaire15]



Difficult areas

- ▶ **Hair & fur**
 - ▶ OIT, overdraw, LOD, quad overshading, deep shadows
- ▶ **Foliage**
 - ▶ OIT, overdraw, LOD, geometry throughput,
 - ▶ Lighting, translucency, AO
- ▶ **Fluids**
 - ▶ LOD & scalability, simulation, overall rendering
- ▶ **VFX**
 - ▶ Need volumetric representation & lighting
 - ▶ Related to [Hillaire15]



Difficult areas

- ▶ Hair & fur
 - ▶ OIT, overdraw, LOD, quad overshading, deep shadows

- ▶ Foliage
 - ▶ OIT, overdraw, LOD, geometry throughput,
 - ▶ Lighting, translucency, AO

- ▶ Fluids
 - ▶ LOD & scalability, simulation, overall rendering

- ▶ VFX
 - ▶ Need volumetric representation & lighting
 - ▶ Related to [Hillaire15]



Pompeii movie

Difficult areas (cont.)

- ▶ **Correct shadows** on everything
 - ▶ Including area lights & shadows!
 - ▶ Extra important with PBR to prevent leakage
 - ▶ Geometry throughput, CPU overhead, filtering, LOD
- ▶ **Reflections**
 - ▶ Hodgepodge of techniques today
 - ▶ Occlusion of specular critical
 - ▶ See Mirror's Edge talk [Johansson15]
- ▶ **Antialiasing**
 - ▶ See [Salvi15] next



Quality challenges

- ▶ Getting the last 5-10% quality can be **very expensive**
 - ▶ While covering a relatively small portion of the screen
 - ▶ Example: **hair & fur** rendering
 - ▶ Improving GPUs in some of these areas **may not benefit** “ordinary” rendering
- ▶ How to build truly **scalable** solutions
 - ▶ Example: Rendering, lighting and shadowing a full **forest**
 - ▶ **Level-of-detail** is a key challenge for most techniques to make them practical

Scalable solutions – screen-space

- ▶ **Sub-surface scattering** went from texture- to screen-space
 - ▶ Orders of magnitude faster
 - ▶ Implicitly scalable + no per-object tracking
 - ▶ Not perfect, but made it practical & mainstream
- ▶ **Volumetric rendering** to view frustum 3d texture
 - ▶ Froxels! See [Wronski14] and [Hillaire15]

Scalable solutions – screen-space

- ▶ Can one **extend screen-space** techniques further?
- ▶ Render multiple **depth layers** to solve occlusion
 - ▶ Multi-layer deep gbuffers [Mara14]
- ▶ Render cubemap to reach **outside of frustum**
 - ▶ Render lower resolution separate cubemap, slow
 - ▶ Render main view as cubemap with variable resolution?
 - ▶ Single geometry pass
 - ▶ Also for fovated rendering

Scalable solutions – pre-compute

- ▶ Traditionally a **strong cut off** between pre-computed & runtime solutions
- ▶ Believe this is going away more – techniques and systems have to scale & cover more of the spectrum:
 - ▶ **Offline** pre-compute: Highest-quality
 - ▶ **Load-time** pre-compute: High-quality
 - ▶ **Background** compute: Medium-quality
 - ▶ **Runtime**
- ▶ Want **flexible tradeoffs** depending on contexts
 - ▶ Artist live editing lighting
 - ▶ Gamer customizing in-game content
 - ▶ Background gameplay changes to the game environment

Scalable solutions – hierarchical geometry

- ▶ Want to avoid wasteful **brute force** geometry rendering
- ▶ **Do your own culling**, occlusion & LOD directly on the GPU
 - ▶ Finer granularity than CPU code
 - ▶ Engine can have more context and own spatial data structures
 - ▶ Combined with GPU information (for example HiZ)
 - ▶ Opportunities to extend the GPU pipeline?
- ▶ **Compute as frontend** for graphics pipeline to accelerate
 - ▶ Avoid writing geometry out to memory
 - ▶ Good fit with procedural geometry systems as well

Takeaways

- ▶ We've gotten very far in the last few years!
 - ▶ Big **transitions**: PBR, Gen4, Compute, explicit APIs
- ▶ We are at the cusp of a **beautiful** future!
- ▶ Build your own **rendering pipelines** & **data structures**
 - ▶ But which ones? All of them! 😊
- ▶ Need reduce **coupling** & further evolve GPU **execution models**

Thanks to everyone who provided feedback!

- ▶ Sébastien Hillaire (@sebhillaire)
- ▶ Christina Coffin (@christinacoffin)
- ▶ John White (@zedcull)
- ▶ Aaron Lefohn (@aaronlefohn)
- ▶ Colin Barré-Brisebois (@zigguratvertigo)
- ▶ Sébastián Lagarde (@seblagarde)
- ▶ Tomasz Stachowiak (@h3r2tic)
- ▶ Andrew Lauritzen (@andrewlauritzen)
- ▶ Jasper Bekkers (@jasperbekkers)
- ▶ Yuiry O'Donnell (@yuriyodonnell)
- ▶ Kenneth Brown
- ▶ Natasha Tatarchuk (@mirror2mask)
- ▶ Angelo Pesce (@kenpex)
- ▶ David Reinig (@d13_dreinig)
- ▶ Promit Roy (@promit_roy)
- ▶ Rich Forster (@dickyjimforster)
- ▶ Niklas Nummelin (@niklasnummelin)
- ▶ Tobias Berghoff (@tobiasberghoff)
- ▶ Morgan McGuire (@casualeffects)
- ▶ Tom Forsyth (@tom_forsyth)
- ▶ Eric Smolikowski (@esmolikowski)
- ▶ Nathan Reed (@reedbeta)
- ▶ Christer Ericson (@christerericson)
- ▶ Daniel Collin (@daniel_collin)
- ▶ Matias Goldberg (@matiasgoldberg)
- ▶ Arne Schober (@khipu_kamayua)
- ▶ Dan Olson (@olson_dan)
- ▶ Joshua Barczak (@joshuabarczak)
- ▶ Bart Wronski (@bartwronsk)
- ▶ Krzysztof Narkowicz (@knarkowicz)
- ▶ Julien Guertault (@zavie)
- ▶ Sander van Rossen (@logicalerror)
- ▶ Lucas Hardi (@lhardi)
- ▶ Tim Foley (@tangentvector)

Questions?



Email: johan@ea.com
Web: <http://frostbite.com>
Twitter: [@repi](https://twitter.com/repi)

References

- ▶ [Andersson07] [Terrain rendering in Frostbite using Procedural Shader Splatting](#)
- ▶ [Hillaire15] [Physically Based and Unified Volumetric Rendering in Frostbite](#)
- ▶ [Wronski14] [Volumetric fog: Unified, compute shader based solution to atmospheric scattering](#)
- ▶ [Salvi15] [Anti-Aliasing: Are We There Yet?](#)
- ▶ [Mara14] [Fast Global Illumination Approximations on Deep G-Buffer](#)
- ▶ [Johansson15] Leap of Faith: The World of Mirror's Edge Catalyst