

GameDevelopers
Conference

MARCH 20-24
SAN JOSE, CALIFORNIA

WHAT'S NEXT
.....GDC:06

www.gdconf.com

GAME DEVELOPERS CHOICE AWARDS

INDEPENDENT GAMES FESTIVAL

GDC MOBILE

SERIOUS GAMES SUMMIT

GAME CONNECTION



CMP

CREATIVE MEDIA PRODUCTION

Extremely Practical Shadows

Tom Forsyth

Programmer, RAD Games Tools

www.radgametools.com

www.eelpi.gotdns.org

Background

- ④ Day job is “Granny” at RAD Game Tools
 - ④ Animation middleware
 - ④ This talk is nothing to do with that
- ④ Previously a “real” game developer
 - ④ Specialised in graphics
 - ④ Particularly shadows and lighting
- ④ This is the sequel to my GDC 2004 talk
 - ④ Lots more depth
 - ④ Bugs fixed!

Volume shadows suck

- ⌘ Limits geometry – must be watertight
- ⌘ Alpha test & pixel kill don't work
- ⌘ High GPU/CPU cost to compute volumes
- ⌘ Hard-edged
- ⌘ Poor scaling with hardware trends
 - ⌘ Fillrate increasing slower than computation
- ⌘ My assertion: “Not The Future”
 - ⌘ But prediction is hard 😊

Shadowbuffers suck too

- ⌘ Spatial aliasing
 - ⌘ Insufficient resolution
 - ⌘ Non-uniform projection
- ⌘ Depth aliasing
 - ⌘ Surface acne
 - ⌘ Peter-Pan syndrome
- ⌘ FOV limits
 - ⌘ Fundamental 180 degree limit
 - ⌘ ~120 degrees in practice

The projection problem

- ④ Shadowbuffer is rendered from light POV
- ④ Then projected into camera POV
- ④ The two do not agree
 - ④ Can violently disagree – “duelling frustums”
- ④ Too many texels in some places
 - ④ Inefficient use of memory and fillrate
- ④ Too few texels in others
 - ④ Visible aliasing

Projection solutions

- ④ Use smarter projections
 - ④ Lots of freedom in the light-space projection
 - ④ Perspective Shadow Maps (PSM)
 - ④ Light-space Persp. Shadow Maps (LiSPSM)
 - ④ Trapezoidal Shadow Maps (TSM)
- ④ All help with spatial aliasing
- ④ Only minor help with depth aliasing
 - ④ Some make it worse!

Perspective Shadow Maps

- ④ Stamminger and Drettakis, 2002
- ④ Performs viewer perspective on scene
 - ④ Objects close to camera get bigger
- ④ Renders shadowbuffer from light POV
 - ④ So objects close to camera get more texels
- ④ Good redistribution of resolution
- ④ Simple to implement basic algorithm

Perspective Shadow Maps

- ⌚ Light rays no longer parallel
 - ⌚ Directional lights become point lights
 - ⌚ Points become directional or “inverse” points
 - ⌚ Lights can cross div-by-zero plane
 - ⌚ Shadow casters can too!
- ⌚ Very difficult to make robust
 - ⌚ Lots of special cases & discontinuities
 - ⌚ Solutions are unintuitive and fragile
 - ⌚ “Friends don’t let friends implement PSM”

Perspective Shadow Maps

 Demo

Light-Space PSM

- ④ Wimmer, Scherzer, Purgathofer 2004
- ④ Similar perspective distortion to PSM
 - ④ But only perpendicular to light rays
 - ④ So light rays are still parallel
 - ④ No discontinuities to worry about
- ④ Complex to implement
 - ④ Needs volume/volume intersection code
 - ④ But there's recipes to follow
 - ④ And you can approximate

Light-Space PSM

 Demo

Trapezoidal Shadow Maps

- ④ Martin and Tan, 2004
- ④ Also distorts light space
 - ④ Perpendicular to light direction like LiSPSM
- ④ Visible frustum approximated by a trapezoid
- ④ Trapezoid is distorted to square shadowbuffer
 - ④ Not directly related to viewer perspective
- ④ Over-distorts in many cases
 - ④ Too much near detail
 - ④ So they use multiple affine projections
- ④ Can maintain some frame-to-frame continuity

Trapezoidal Shadow Maps

- ④ Distortion is segmented, not affine
 - ④ Solved by pixel-shader lookup (1D texture)
 - ④ But depth is still tricky
- ④ Finding the trapezoid is complex
 - ④ Problems with “infinite” view frustum
 - ④ Iterative process finds distortion segments
- ④ Designed for large ground planes
 - ④ Not very good at more “3D” worlds
 - ④ Or odd-shaped ones (see later!)

Trapezoidal Shadow Maps

 Demo

Smart projection problems

- ④ None of them solve duelling frustum case
 - ④ Camera looking at light source
 - ④ Two frustums are directly opposed
 - ④ Projection has no degrees of freedom
 - ④ All do the same as standard projection
- ④ None of them solve omni lights
 - ④ Cannot render FOV >180 degrees
 - ④ Guaranteed duelling frustum case!

What I tried

④ Simplified LiSPSM

- ④ Affine perspective distortion
- ④ Perpendicular to light rays
- ④ In direction of viewer Z

④ I know more about my scene

- ④ Vanilla algorithm assumes near-infinite view
- ④ I have more control over what is in frustum
- ④ So I can use a simpler algo
(see later for details)

Multi-frustum partitioning

- ④ Splits scene into multiple light frustums
 - ④ Each frustum rendered separately
- ④ Solves the two big problems
 - ④ Duelling frustums
 - ④ Omni lights
- ④ Helps in other ways
 - ④ Copes gracefully when smart projection fails
 - ④ More control over frustum contents
 - ④ ...allows “dumber” smart projection

MFP assumptions

- ⌘ Each light has multiple frustums
- ⌘ Each object is caster and/or receiver
- ⌘ Receivers use only one frustum
 - ⌘ Only care about receivers visible to viewer
- ⌘ Casters can affect multiple frustums
 - ⌘ Still care about casters not visible to viewer
- ⌘ Algorithm mainly concerns receivers
 - ⌘ Casters are easy to deal with
 - ⌘ Standard GPU clipping does the right thing

MFP basic algorithm pt1

- ④ For each receiver in scene
 - ④ Find screen size of bounding volume
 - ④ Texels-per-screen-pixel global quality setting
 - ④ Result is texels-per-world-meter value
- ④ Each light frustum has
 - ④ Frustum direction/FOV
 - ④ I use a circular cone for simplicity
 - ④ Texels-per-degree variable
 - ④ List of receivers that use frustum

MFP basic algorithm pt2

- ④ Receiver calculates texel density for light
 - ④ (ignore receivers not visible to viewer or light)
 - ④ Find FOV angle relative to light position
 - ④ So find texels per degree of light angle
- ④ Receiver looks for suitable frustum
 - ④ Must contain receiver volume in frustum
 - ④ Must have at least enough texels-per-degree
 - ④ But not too high - wastes fillrate & memory
 - ④ Heuristic factor of $\sim 4x$ allowed

MFP basic algorithm pt3

- ④ If no frustum found, try to enlarge one
- ④ Increase frustum FOV
 - ④ ...and alter direction of frustum
 - ④ Union of new receiver and existing list
 - ④ Texture size limit (texels per degree * FOV)
 - ④ FOV distortion limit – around 120 degrees
- ④ Increase texels per degree
 - ④ Not past factor of 4x to any receiver in list
 - ④ Prevents fillrate & memory waste

MFP basic algorithm pt4

- ④ If still no suitable frustum, make new one
 - ④ Exactly fits this receiver
 - ④ Exactly matches texels-per-degree
 - ④ Can be expanded to include other receivers
- ④ Receiver may exceed single frustum
 - ④ Too large to fit in 120 degree FOV
 - ④ Too large for texture size restrictions
 - ④ May need chopping
(see later for more)

MFP + smart projections

- ④ MFP can use any projection
- ④ Different projection for each frustum
 - ④ Can use PSM where no singularity problems
 - ④ Use TSM/LiSPSM/standard elsewhere
- ④ Each frustum is well-constrained
 - ④ Can bias away from PSM singularities
 - ④ TSM not good at large/odd-shaped frustums
 - ④ LiSPSM gets more information for fine-tuning
- ④ Guaranteed fallback to standard proj.

MFP results

- ⌚ Consistent texel density
 - ⌚ 4x best/worst ratio – filtering can cope
 - ⌚ Objects close to viewer get more texels
 - ⌚ Arbitrarily large scale variation in scene
- ⌚ Distance from light is accounted for
 - ⌚ Not perfect – surface may be edge-on
- ⌚ Duelling frustums solved by partitioning
 - ⌚ Usually only need 3-4 regions at extreme
 - ⌚ Smart projections help, but not required

MFP results

- ④ <120 degree FOV chunks
 - ④ Wide/omni lights use multiple frustums
- ④ Better than cube maps
 - ④ “Cube sides” can be different resolutions
 - ④ Unused/invisible “sides” never considered
 - ④ Omnis with nothing above them but sky
 - ④ Omnis beside viewer, shining into scene
- ④ Omnis always have a duelling frustum
 - ④ “Side” with duelling frustum is partitioned

MFP practical results

- ④ Retrofitted to StarTopia (2001)
 - ④ RTS/"god game"
 - ④ Player-built world – no preprocessing
- ④ 2D grid-aligned world
 - ④ But shadowing system doesn't know this
 - ④ There's no cheating!
 - ④ Many objects larger than a grid square
 - ④ Many non-aligned characters
 - ④ World is bent into a ring

MFP practical results

- ⌚ No gameplay or artwork changed
 - ⌚ Alpha-test & clip planes used extensively
 - ⌚ Already had lighting info, but no shadows
 - ⌚ Many objects contain lights
 - ⌚ All omnis – rampant duelling frustums!
- ⌚ Very few lights altered
 - ⌚ Some moved to nicer-looking places (lamps)
 - ⌚ Shadows for “mood lights” turned off
- ⌚ This implementation is slow
 - ⌚ Engine is bad at multiple rendering passes

MFP results

🕒 Demo

- 🕒 “Patch” downloadable from
www.eelpi.gotdns.org

(StarTopia is ~\$3 from Amazon)

- 🕒 Email me for debug modes shown here

MFP problems - chopping

- ④ Large/close receivers are problems
 - ④ No single frustum can cover whole object
 - ④ But first reject against light & viewer volumes
 - ④ Problem part may be out of range/view
- ④ May already have finer chunks
 - ④ Material / bone limit boundaries
 - ④ Procedural, e.g. landscape tiles
 - ④ Precalculated split points (long walls & floors)
- ④ But sometimes they still need chopping

MFP problems - chopping

- ⌚ Not very common
 - ⌚ So pick simplest method
- ⌚ Split into arbitrary pieces
 - ⌚ Need at worst six: cube-map chop planes
- ⌚ Find frustum/SB for each piece
- ⌚ Render object with all shadowbuffers
 - ⌚ Ensure out-of-frustum reads = “dark”
 - ⌚ Take maximum brightness of all results
 - ⌚ Hardware limits may require multiple passes

MFP problems - popping

- ⌚ Sudden aliasing change between frames
 - ⌚ Object moves from frustum to frustum
 - ⌚ Frustum changes direction/FOV
 - ⌚ Frustum changes resolution
- ⌚ Frustums change as objects move
 - ⌚ Also as viewer or light moves
 - ⌚ ...but popping is far less visible then
- ⌚ Frustum algorithm is greedy
 - ⌚ Small movements can cause large changes
 - ⌚ Changes are not localised to moving objects

MFP problems - popping

- ④ Limit changes to existing frustums
 - ④ Texel->world mapping cannot change
 - ④ Can pan by whole texels
 - ④ Can change texture size within limits
 - ④ Can change scissor to avoid extra overdraw
- ④ Make new frustums “steal” from old ones
 - ④ Old ones removed when no objects in them
 - ④ Limits number of old “fragmented” frustums

MFP problems - popping

- ④ Blend when object changes frustum
 - ④ Keep adding to both old and new
 - ④ Render twice & crossfade
- ④ But object may not fit in old frustum
 - ④ New position, size, etc
 - ④ There's a reason it changed frustum!
 - ④ Will still pop, but it's rare
- ④ Pop not as visible on animated objects
 - ④ Only need this for moving non-anim, e.g. cars

MFP problems - seams

- ④ Adjacent objects in different frustums
 - ④ Slightly different aliasing between the two
- ④ Only a problem with smooth joins
 - ④ Floor/land tiles, wall sections
 - ④ Hidden by noisy textures
- ④ Discontinuity invisible over sharp joins
- ④ Soft shadows help hide seams
 - ④ Screen-space blur effective, but expensive
 - ④ ATI “Parthenon” demo explores solutions

MFP problems - seams

- ⌚ Store smooth connectivity between receivers
- ⌚ Grow object volumes to include neighbours
 - ⌚ Frustum can now cover both sides of seams
 - ⌚ Does not create significantly larger frustums
- ⌚ Both objects on a seam rendered twice
 - ⌚ Once with each frustum
 - ⌚ Blend at 50:50 at seam edge
- ⌚ Requires vertex weights
 - ⌚ Annoying preprocessing for general meshes
 - ⌚ Easy with procedural geometry (floor tiles)

Render target management

- ④ Allocate a few large textures
- ④ Quadtree allocation inside each
 - ④ Very fast, low fragmentation, easy to write
- ④ Each allocation uses same texture as last
 - ④ Minimise render-target changes
 - ④ Only switch when you have to
 - ④ Then stay switched
- ④ Leave 1-textel borders
 - ④ Careful with filtering, can't use CLAMP mode

Depth aliasing

- ⌚ Lots of biases to use and tweak
 - ⌚ 24 bits is still not enough
 - ⌚ Low precision far from light, but may be near viewer!
- ⌚ Bias too high
 - ⌚ Shadows detach from feet – “Peter-Pan” syndrome
- ⌚ Bias too low
 - ⌚ Incorrectly self-shadows - “surface acne”
 - ⌚ Render backfaces – terrible Peter-Panning!
 - ⌚ Midpoint rendering – expensive, tricky
- ⌚ No one bias value works for an entire scene

ID buffers

- ④ Store object ID (not depth) in shadowbuffer
- ④ Compare to ID of rendered object
- ④ If they match, object is lit
- ④ No math
 - ④ No biases, no epsilons, no precision problems
- ④ 256 IDs usually enough
 - ④ Chances of collision are low
 - ④ Collision causes no shadow, not extra shadows
- ④ 65536 IDs enough
 - ④ You have how many objects in your scene?

ID buffers

⌚ Dark halos

- ⌚ Sampling at object edge picks up IDs behind
- ⌚ Can use “priority buffer” – requires ordering

⌚ Point-sample nearest four texels

- ⌚ In shadow if no ID matches
- ⌚ Shadows shrink by a texel – not noticeable

⌚ No self-shadowing

- ⌚ Could use ID per face
- ⌚ ...needs lots of IDs & hardware support

Depth + ID buffers

- ④ ID buffers used for inter-object shadows
- ④ Depth used for self-shadowing

```
if ( buf.ID != obj.ID )
```

```
    in_shadow();
```

```
else if ( buf.depth < obj.depth )
```

```
    in_shadow();
```

```
else
```

```
    draw_lit();
```

Depth + ID buffers

- ④ Depth range only spans a single object
 - ④ Inter-object shadows done with IDs
- ④ Few bits needed
 - ④ Blade 2 on Xbox1 used only 8 bits
 - ④ ...one depth buffer per character
 - ④ 8 bits over 2 meters = ~1cm resolution
 - ④ Barely enough – 16 bits probably wiser
- ④ Bias can be set differently for each object
 - ④ Set once per object, no scene dependencies

Depth + ID buffers

- ④ 8 bit ID + 8 bit depth = 16 bits
 - ④ But probably need more depth bits
- ④ Some effects need a global depth
 - ④ Depth of field
 - ④ Soft-edged shadows
 - ④ All are more tolerant of precision/bias errors
- ④ Can just use regular 24-bit Z buffer
 - ④ 8-bit ID stored in stencil channel
 - ④ If hardware supports it
 - ④ Still use per-object bias

Depth + ID buffers

- ④ Demo
- ④ 8-bit ID - ID 0 reserved for floor tiles
- ④ Other 255 IDs assigned in render order
 - ④ Front to back (be nice to your Z buffer!)
 - ④ Chance of collisions low – widely spaced
- ④ Depth bias scaled by object radius
 - ④ Could have artists fine-tune them per-object
 - ④ But I have no artists 😊

Soft shadows

- ⌚ PCF looks bad
 - ⌚ Can get away with it on noisy surfaces
 - ⌚ More taps are expensive, little extra quality
 - ⌚ Not depth-dependent (blurry feet look wrong)
- ⌚ Smoothies get good results
 - ⌚ Chan and Durand, 2003
 - ⌚ But needs geometric volumes!
- ⌚ Willem de Boer's variant of smoothies
 - ⌚ Entirely image-based

Questions...

tom.forsyth@eelpi.gotdns.org

www.eelpi.gotdns.org

References

- ④ “Perspective Shadow Maps”, Stamminger and Drettakis, Proceedings of ACM SIGGRAPH 2002
- ④ “Light Space Perspective Shadow Maps”, Wimmer, Scherzer, Purgathofer, Eurographics Symposium on Rendering 2004

References

- ④ “Anti-aliasing and Continuity with Trapezoidal Shadow Maps”, Martin and Tan, Eurographics Symposium on Rendering 2004
- ④ “Combined Depth and ID-Based Shadow Buffers”, Kurt Pelzer, Game Programming Gems 4, Charles River Media 2004

References

- ④ “Rendering Fake Soft Shadows with Smoothies”, Chan and Durand, Proceedings of the Eurographics Symposium on Rendering 2003
- ④ “Smooth Penumbra Transitions with Shadow Maps”, Willem de Boer, www.whdeboer.com

References

- ④ “StarTopia”, developed by Muckyfoot Productions, published by Eidos Interactive, 2001