



Squeezing Performance out of your Game with ATI Developer Performance Tools and Optimization Techniques

Jonathan Zarge, Team Lead Performance Tools
Richard Huddy, European Developer Relations Manager
ATI Technologies, Inc.



Outline

- DX9 Optimization Techniques: Richard Huddy – 25 minutes
 - Graphics pipeline overview
 - Optimizing at each pipeline stage
 - Tips for writing efficient code
- Performance Tools: Jonathan Zarge – 30 minutes
 - ATI developer performance tools overview
 - PerfDash
 - our new real-time performance analysis tool
 - PIX plug-in
 - track ATI hardware counters in PIX
 - gDEDebugger
 - OpenGL debugging - now with support for ATI performance counters
- ATI Content Creation Tools: Jonathan Zarge – 5 minutes



DirectX 9 Optimization Techniques



Some early observations

- Graphics performance problems are both commoner and rarer than you'd think...
- The most common problem is that games are CPU-limited
- But you can certainly think of that as a graphics problem...
 - As it's most often caused by graphics abuse...



There's plenty of mileage in...

- Instancing

- Available on all ATI's recent hardware
- That's all SM3 hardware
- On ATI's SM2b hardware thru a simple backdoor...
- Use Instancing for objects up to ~100 polys

- Batching

- You will be CPU limited if you don't send your triangles in large groups
 - You can think of this as pretty much a fixed overhead per Draw call in DX9, **much less** in DX10...



DirectX9 State Changes

- Top 5 by cost:
 - SetPixelShaderConstant()
 - SetPixelShader()
 - SetVertexShaderConstant()
 - SetVertexShader()
 - SetTexture()
- So try to avoid these when you can



Unified shaders?

- Think cross platform...
 - Xbox 360
- When it happens the dynamics of PC graphics will change radically



Shall we have a target?

- 1600x1200 and 1280x1024 (at least)
- 85Hz (then lower refresh rates will just work)
- 4xAA or 6xAA – so pixels can look good
- Because of the variability of the platform it makes no sense to ask blindly if we are pixel-limited or vertex-limited etc.
 - [And with U.S. that idea stops making sense...]



Down the bottleneck pipeline...

- Cache re-use
 - VFetch, Vertex, texture, Z
 - All caches are totally independent of each other...
- Vertex shaders
- Pixel shaders
- Z buffer
- Frame buffer



The pre-VS cache I

- Is purely a memory cache
- Has a common line size of 256 bits
 - (That's 32 bytes)
- Is accessible by all vertex fetches
- Is why vertex data is best aligned to 32 bytes or 64 bytes
 - 44 is very much worse than 64
 - Roughly sequential access should be your aim



The pre-VS cache II

- Because it's purely a memory cache...
 - Multiple streams can both help and hinder.
 - Multiple streams with random access is doubly bad...
 - Generally expect 0% to 10% hit for using additional streams



Vertex Engines I

- Consider compressing your vertex data if that helps you line things up with the 32 byte cache line...
 - Decompress in the Vertex Shader
 - Store compressed data in VB
- See previous slide for the point...
- This can be a significant win if it achieves some key alignment objectives



Vertex Engines II

- HLSL is your best approach...
 - We recommend that you compile with optimisations disabled
 - we'll get to know more that way and usually do better
- Expect one op per clock per pipe
 - Sometimes you'll get 2 ops instead...
 - Masking out unused channels helps
 - You can get up to 5 ops at once!
- I've never seen a game which is vertex-throughput limited at interesting resolutions on modern hardware



The post-VS cache

- Only accessible when using indexed primitives (can give you 'free' triangles)
- Operates as a FIFO
- Use D3DXOptimizeMesh()
- Is 14 entries for triangles, 15 for lines and 16 for points
- Cache Size is independent of vertex format!
- Use highly local winding for best results
- Flushed between DrawPrim() calls

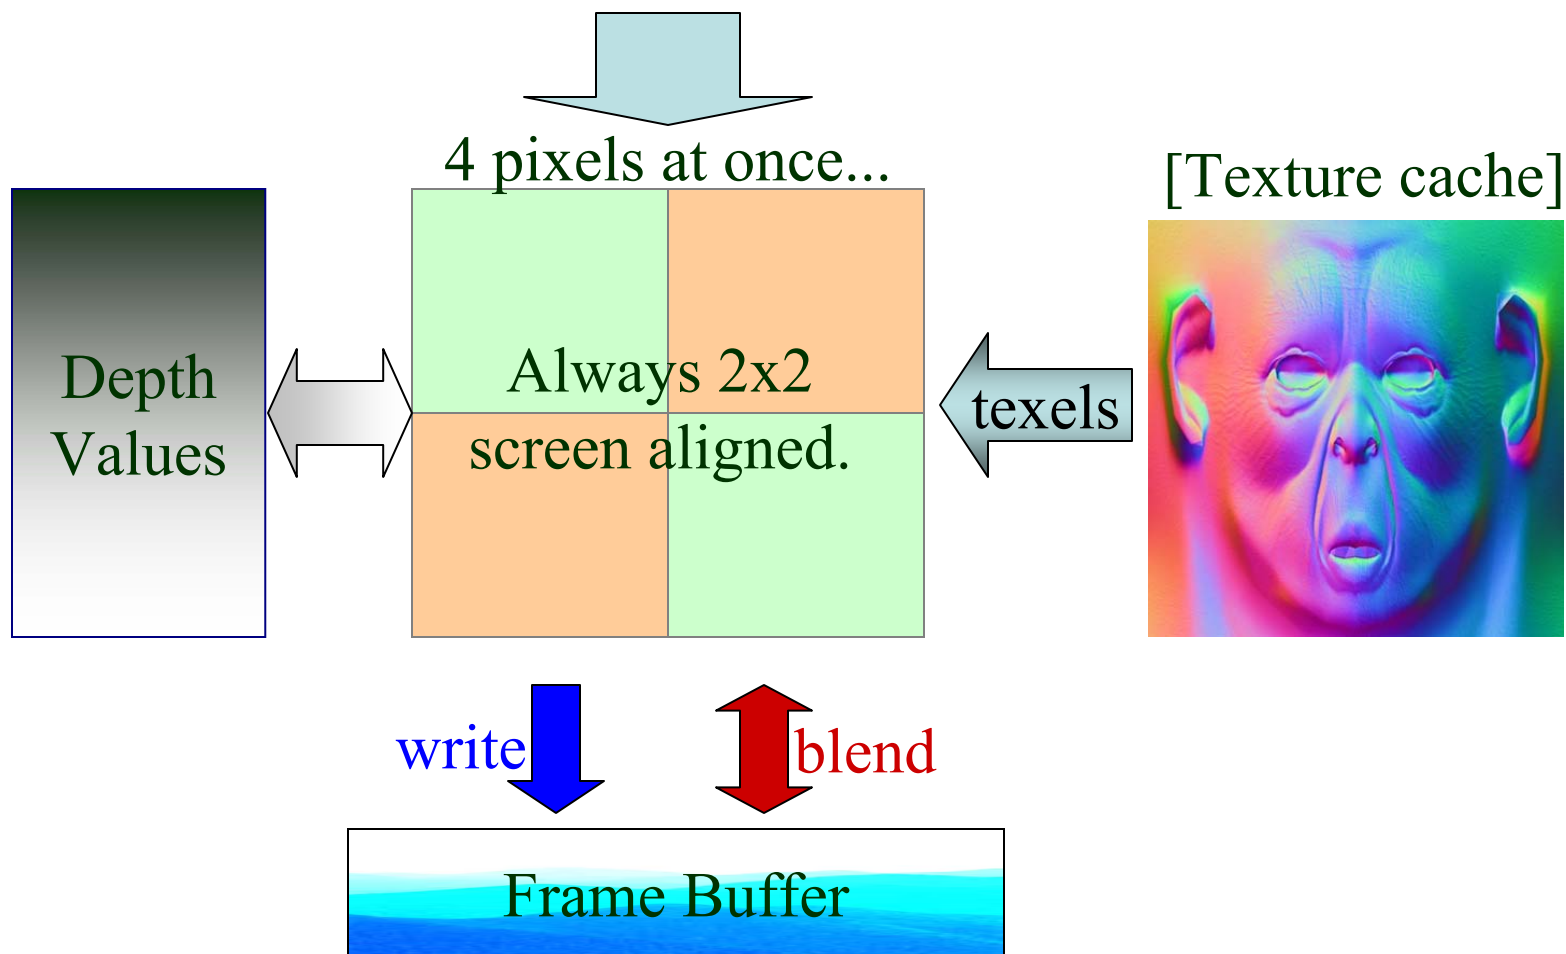


Triangle setup

- Never a bottleneck
- Just joins vertices into triangles
- Feeds the rasterizer which simply hands out quad pixel blocks to draw



A Quad-Pixel Processing Unit





Texture cache

- Probably smaller than you'd think...
 - Unless you thought "only a few KB"
- Partitioned over all active textures
 - So heavy multi-texturing can really hurt
 - Modern hardware has efficient fully associative caches
- Wrecked by random access!
 - Often from bump-map into env-map
 - Needs reuse to show benefits (i.e. don't minify!)
- Usually contains uncompressed data
 - At 8, 16, 32 or more bits per texel
 - Some hardware stores DXT1 in compressed format
- Texture fetches are per-pixel



Making Z work for you...

- We're faster at rejecting than at accepting...
 - So draw roughly front to back
 - For complex scenes consider Z pre-pass (not for `depth_complexity=1!`)
 - Take care to `Clear()` Z (and stencil)
- Although Z is logically at the end of the shader that's not the best way



Making Z work for you...

- Note that ATI hardware can do double speed Z/Stencil only work when:
 - Color-writes disabled
 - AA enabled
 - Good for general rendering
 - AA is your default, yes?
 - That's up to 32 AA Z values per clock



Depth Values

- Can come from:-
 - Actual Z buffer (slow)
 - Compressed Z (fast & lossless)
- Your pixel can be Z-tested away before the shader has run at all!
- If you are performing any Z compare then please try hard not to write to oDepth
- Remember that depth values are per-sample...



Bashing the depth buffer

- You can reduce the huge(*) early Z benefits by...
 - Writing oDepth
 - Kills compressed Z and early Z
 - Using alpha-test etc on visible pixels
 - decompresses Z values
 - Changing the Z compare mode (sometimes)
 - Can disable Hi-Z
 - E.g. from LESS to GREATER

(*) Top class h/w can reject 256 pixels per clock!



The PS Unit I

- Shorter shaders generally faster
 - And we can cache them on chip too...
- At the high end there is roughly 4 times as much ALU power as texture power
- This ratio will only go up
 - Because available bandwidth doesn't rise as fast as chip density
- So generally push more maths into here



The PS Unit II

- Is a 4D vector processor
 - So try to match your math to your needs
 - i.e. Mask out unused channels
- Trust the compilersu to schedule things well:-
 - You don't worry about scheduling...
- PS runs once per pixel...



FB (Fog and) Blend

- Is not part of the PS unit
 - You can think of it as a special function of the memory controller
- Although there are lots of latency hiding tricks here...
 - This is still probably the easiest place to get B/W limited
- So disable blend whenever possible



Pure FB optimisations

- Fewer bits are written faster...
 - $16\text{BPP} > 32\text{BPP} > 64\text{BPP} > 128\text{BPP}$
 - (here '>' means faster)
- Blending is slower than not
 - Often by more than a factor of 2
- ATI: Surfaces are 'faster' when allocated earlier!



PS Dynamic Flow Control

- DFC can be a significant benefit...
 - But only when the selection coherency is at least as big as the hardware batch size

Hardware Batch Size

X1800	16 pixels
X1900	48 pixels
Xenos	48 pixels



Conclusion...

- Several classes of optimisation:
 - Pushing things back up the pipe:
 - E.G. Cull early, not late
 - Getting better parallelism:
 - E.g. Use write masks in your shader code to allow SIMD
 - Doing less is faster than doing more:
 - E.g. Short shaders are faster
 - Understand what is cached:
 - 32 byte vertices are fast! 16 bytes are faster...



Performance Tools

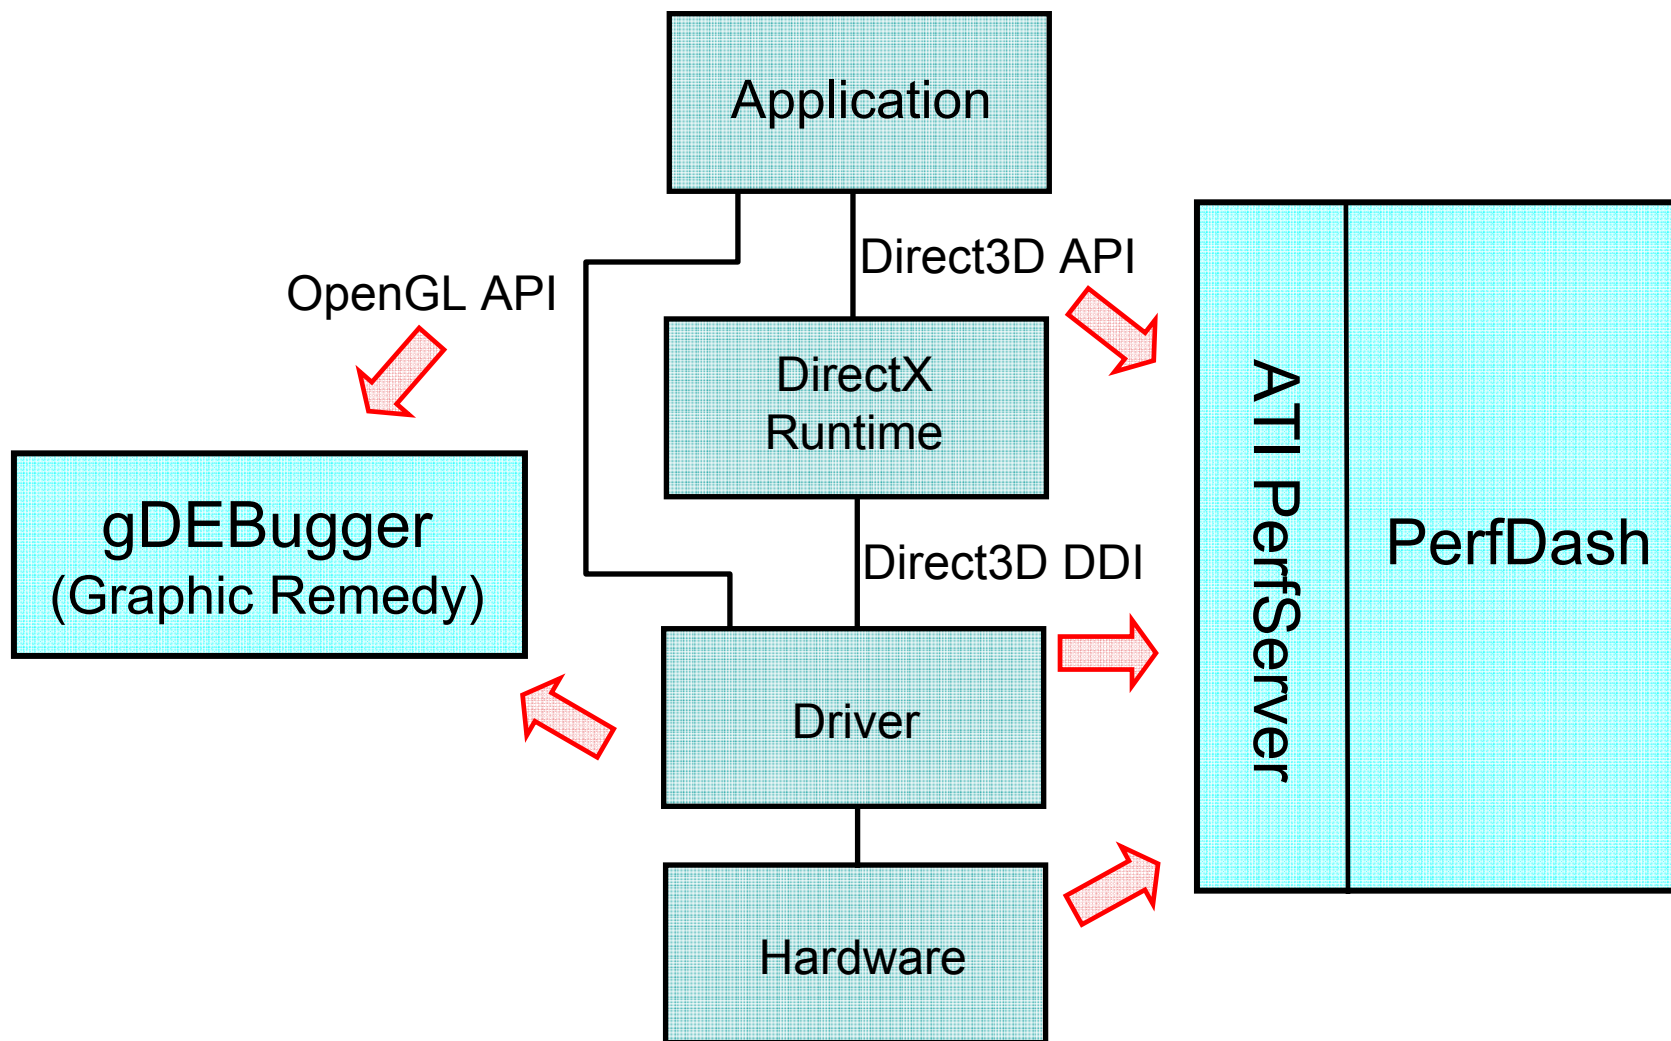


The Plan

- Overview of ATI developer performance tools
- PerfDash
- ATI PIX Plugin
- gDEBugger
- ATI content creation tools



ATI Developer Performance Tools





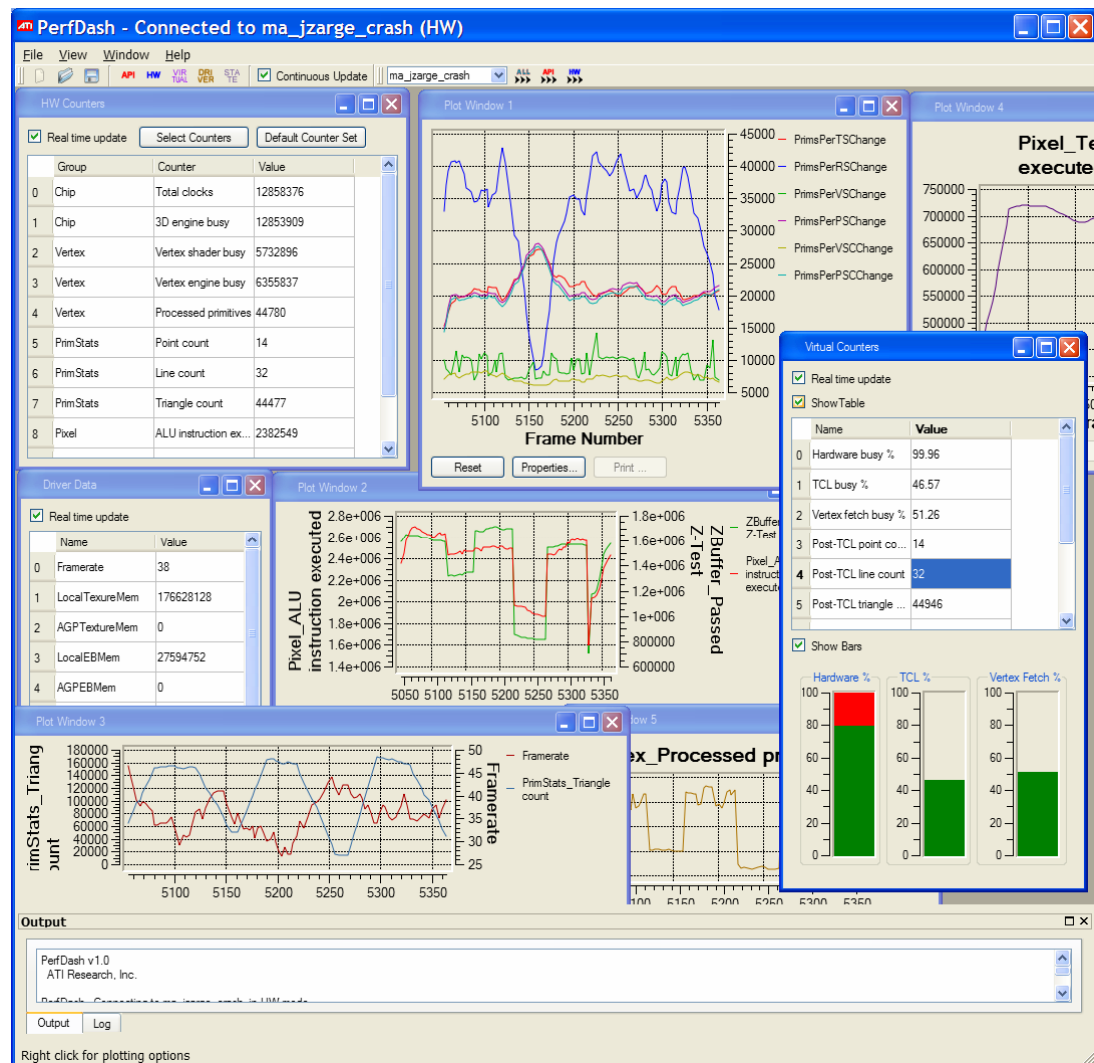
The Plan

- Overview of ATI developer performance tools
- **PerfDash**
- ATI PIX Plugin
- gDEBugger
- ATI content creation tools



PerfDash Overview


- PerfDash: Performance Dashboard
- Real-time visualization:
 - API statistics
 - Hardware counters
 - Driver data
 - "Virtual" counters
- Local or remote performance profiling
- Overriding rendering states
- Loading/saving session data and preferences
- No special driver
- No code modifications
- Plugin architecture





PerfDash Features: Toolbar

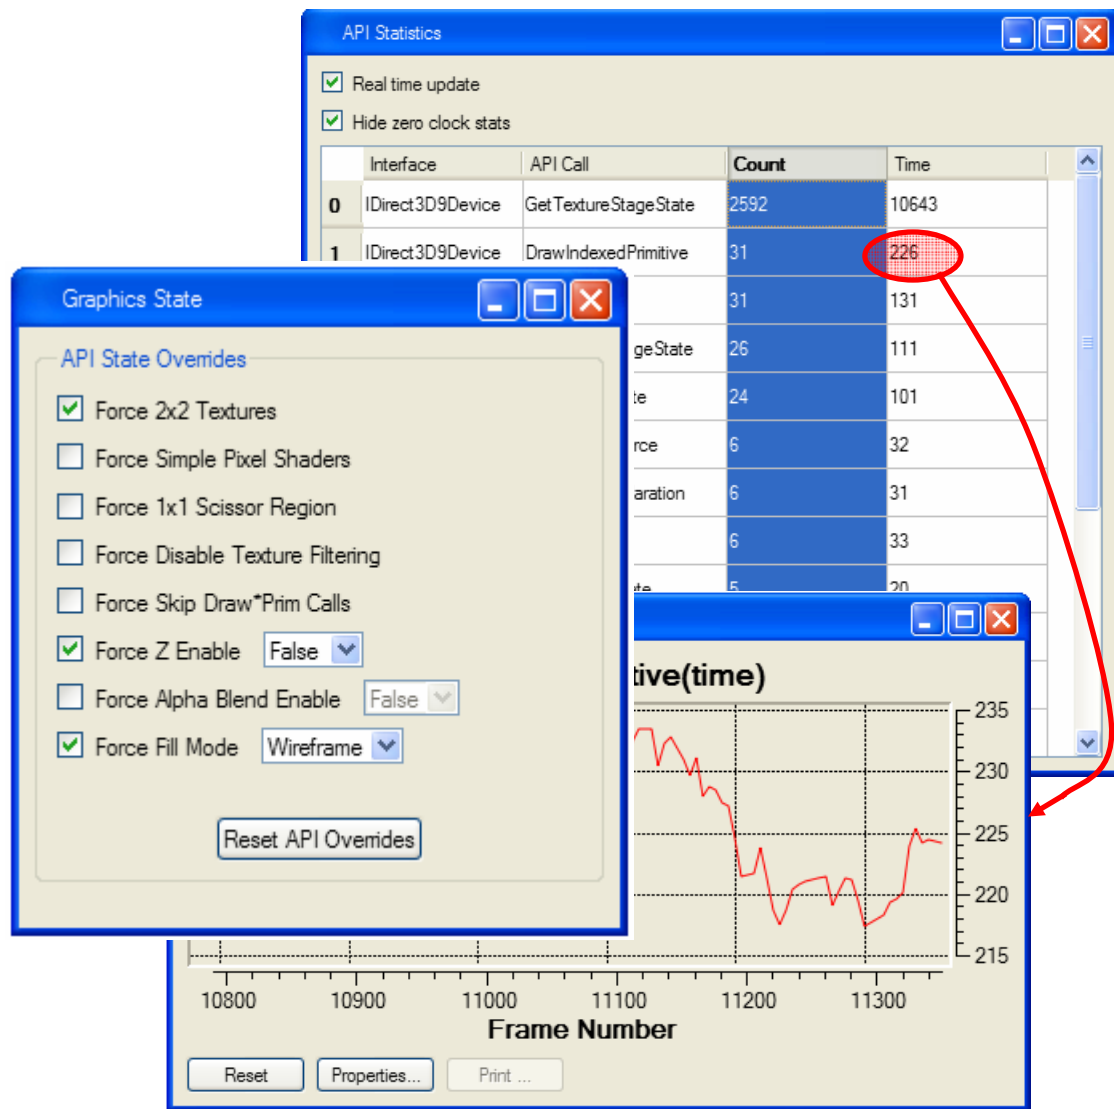
- Toggle global data collection
- Connect to local or remote machine
 - API, Hardware, All modes
 - Performance server must be running on target machine
 - Virtually no performance impact if not running PerfDash





PerfDash Features: API Statistics

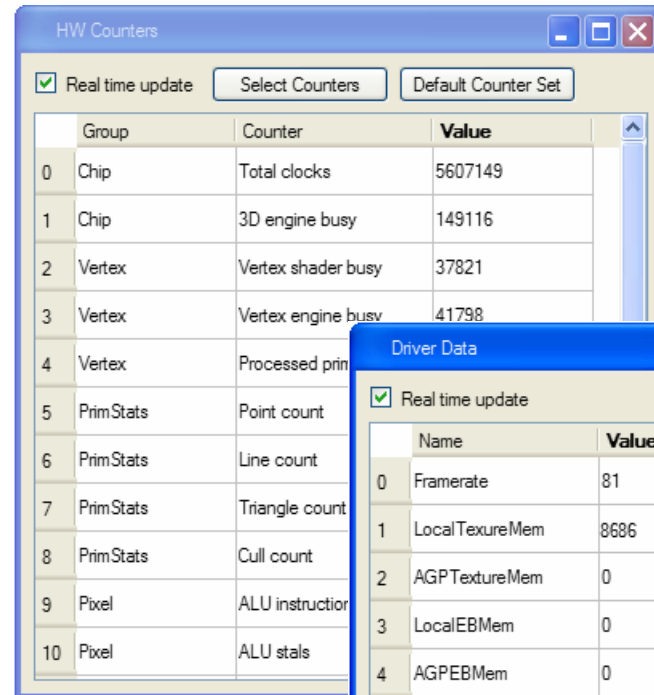
- Per-frame API call data
- Sorting of API call counts and times
- Flexible plotting of all numeric data
- Plot window properties control appearance
- Real-time state overrides





PerfDash Features: Hardware

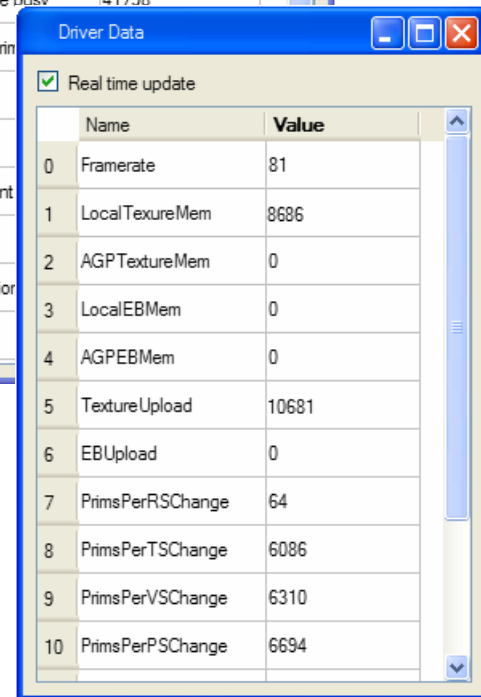
- Hardware counter values
 - 3D/TCL clocks
 - Primitive counts
 - ALU instructions executed
- Select custom set of counters
- Driver data
 - Framerate
 - Memory in use
 - Prims per state change
- Plotting of all numeric data



HW Counters

☒ Real time update Select Counters Default Counter Set

	Group	Counter	Value
0	Chip	Total clocks	5607149
1	Chip	3D engine busy	149116
2	Vertex	Vertex shader busy	37821
3	Vertex	Vertex engine busv	41798
4	Vertex	Processed prim	
5	PrimStats	Point count	
6	PrimStats	Line count	
7	PrimStats	Triangle count	
8	PrimStats	Cull count	
9	Pixel	ALU instruction	
10	Pixel	ALU stals	



Driver Data

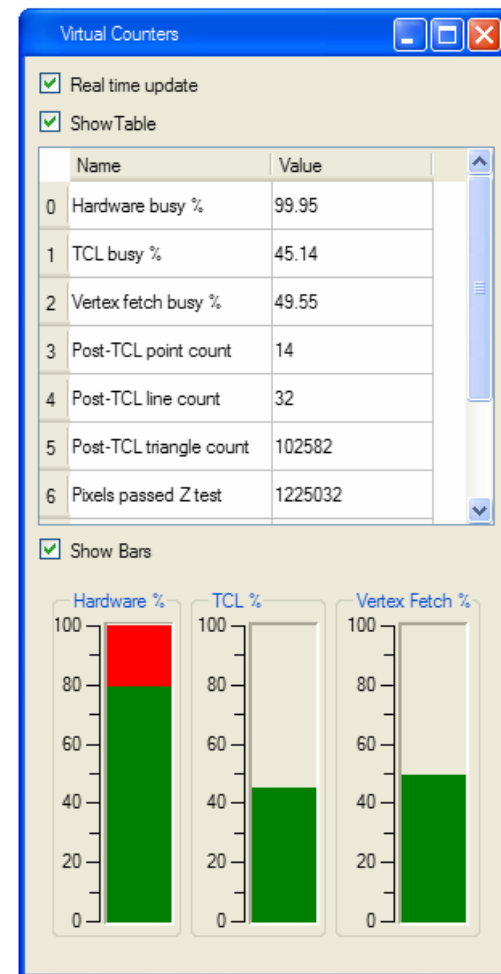

☒ Real time update

	Name	Value
0	Framerate	81
1	LocalTextureMem	8686
2	AGPTextureMem	0
3	LocalEBMem	0
4	AGPEBMem	0
5	TextureUpload	10681
6	EBUpload	0
7	PrimsPerRSChange	64
8	PrimsPerTSChange	6086
9	PrimsPerVSChange	6310
10	PrimsPerPSChange	6694



PerfDash Features: Virtual Counters

- Virtual (derived) counters
 - Hardware busy %
 - TCL busy %
 - Pixels passed z-test
- Temperature bars
- Plotting of all numeric data



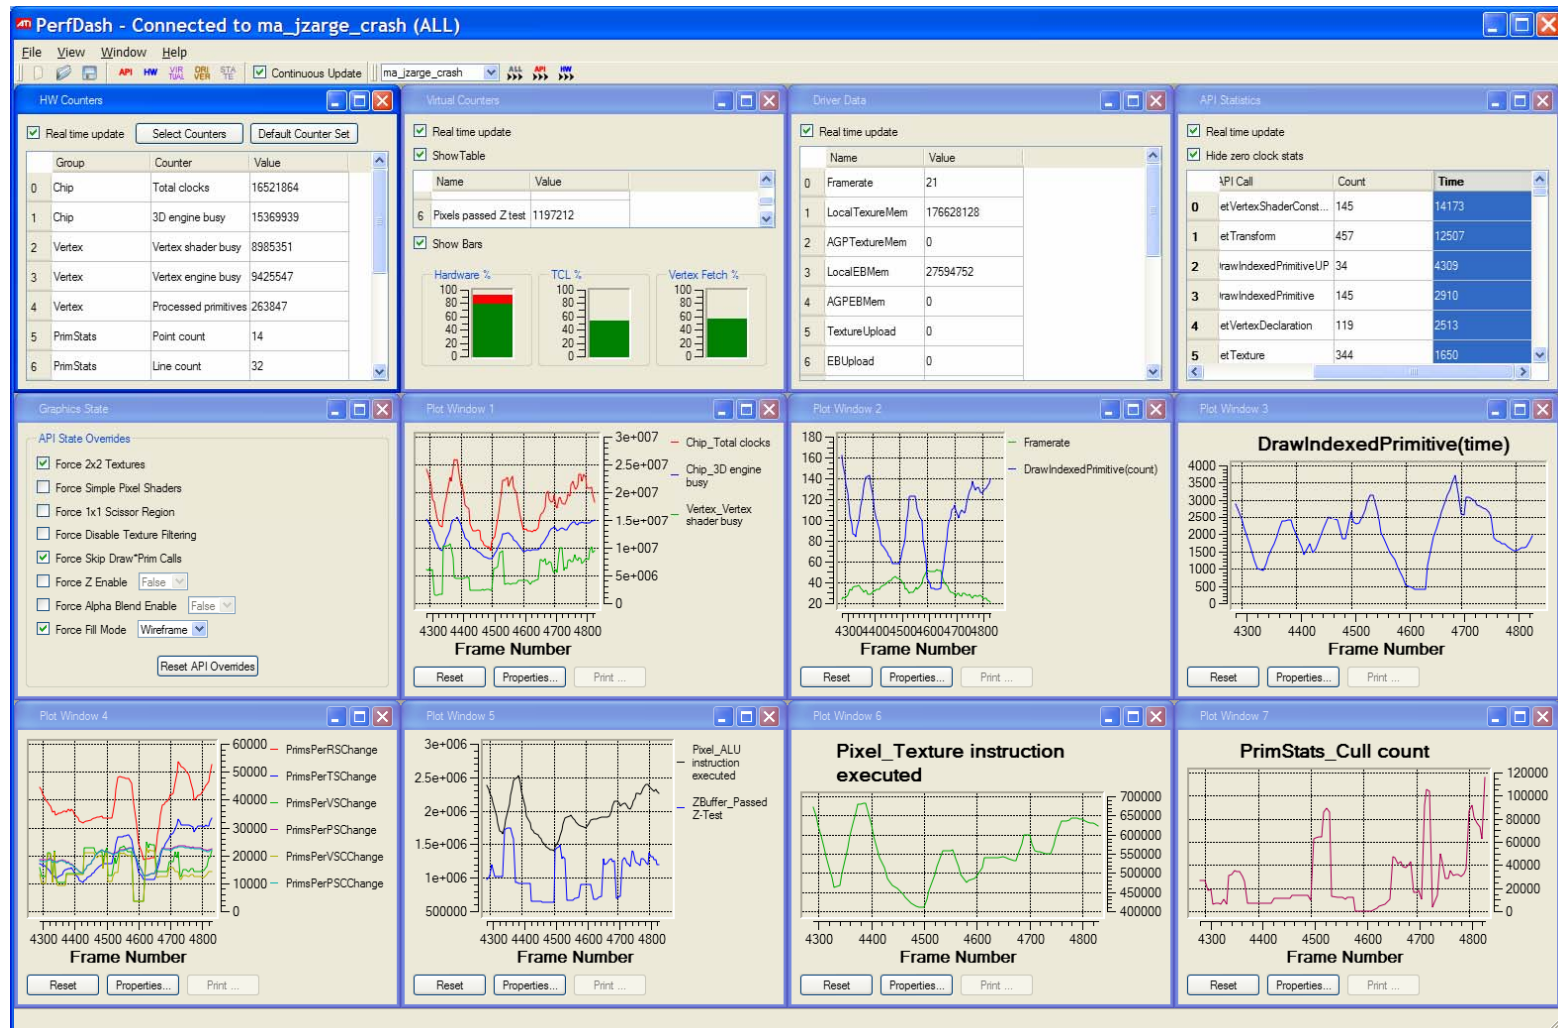


Performance Tuning w/PerfDash

- CPU vs GPU balance
 - Is the GPU saturated?
- TCL percentage
 - Is vertex processing the bottleneck?
- Pixel shader bottleneck
- Draw*Primitive per state change
 - Good indication of batch size
- Inefficient use of API
- High memory usage
- Texture bandwidth and filtering
- Efficient use of z-buffer and sorting
- Wireframe for geometry visualization



PerfDash Demo





PerfDash Future

- Future enhancements
 - Record/playback/step through API calls
 - Support external plugins
 - Support for different platforms
 - More data: pixel stats, render states, buffers, surfaces
 - Bottleneck detection
 - Custom virtual counters
 - Support for OpenGL applications
- PerfDash release schedule
 - Beta in Q1 06
 - First release Q2 06



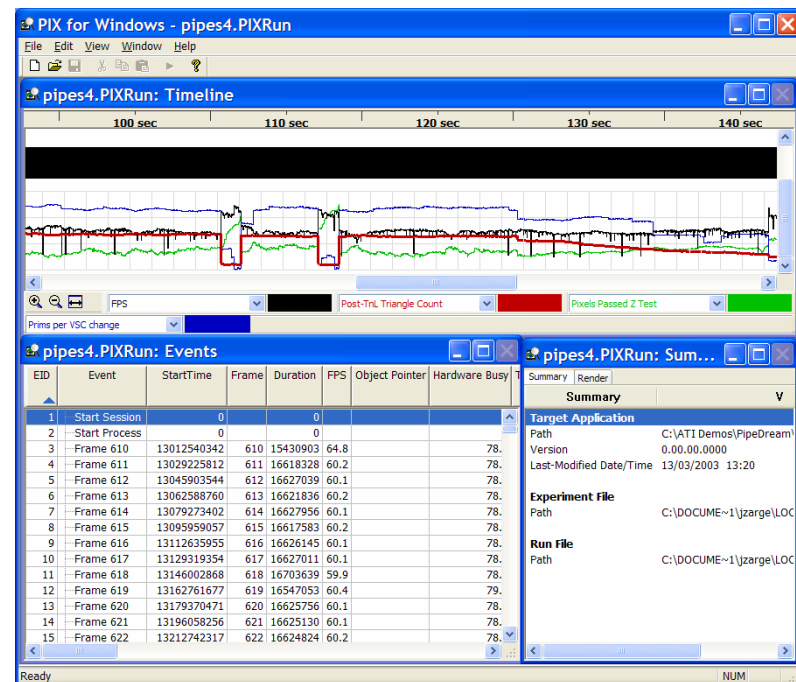
The Plan

- Overview of ATI developer performance tools
- PerfDash
- **ATI PIX Plugin**
- gDEBugger
- ATI content creation tools



PIX for Windows

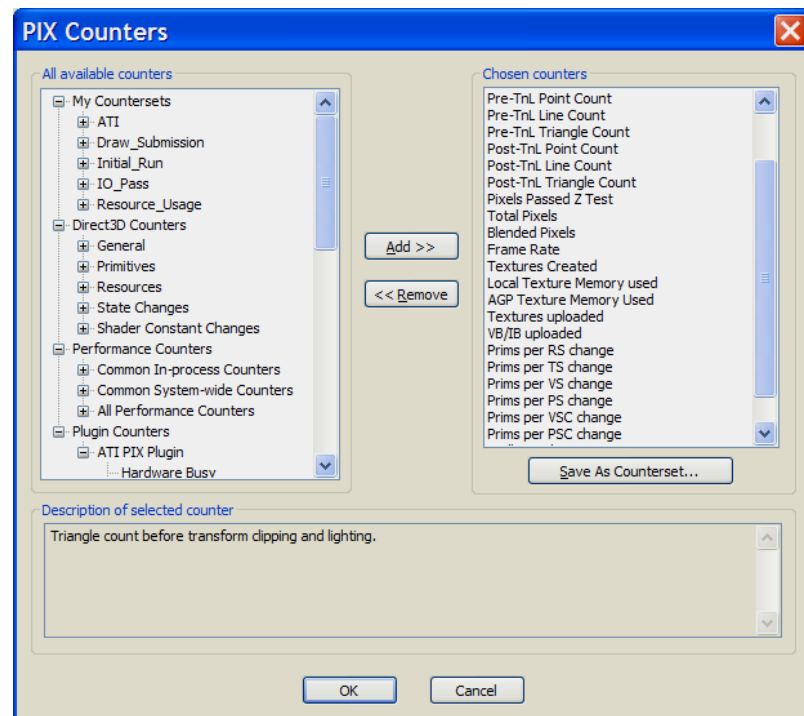
- Based on Xbox PIX
- Record/playback/visualize model
- Direct3D API level
- Numerous included counters
- Plugin API for additional counters
- Image comparison
- Display render states, textures, objects, shaders, vertex decl.
- 3 recording modes
 - Gather statistics (counters)
 - Record D3D calls
 - Record playable stream information





ATI PIX Plugin

- Communicates with ATI Driver (through D3D API)
- Virtual counters
 - Computed from actual hardware counters/driver data
 - Can be graphed with other performance counters
- Counter examples:
 - Hardware/TnL Busy
 - Vertex Fetch Busy
 - Triangle/Line/Point Count
 - Total/Blended/Pass Z Pixels
 - Local/AGP Texture Memory Used
 - Primitives per RS/TS/PSC/VSC
 - Stalls on Flip/VB






The Plan

- Overview of ATI developer performance tools
- PerfDash
- ATI PIX Plugin
- **gDEBugger**
- ATI content creation tools



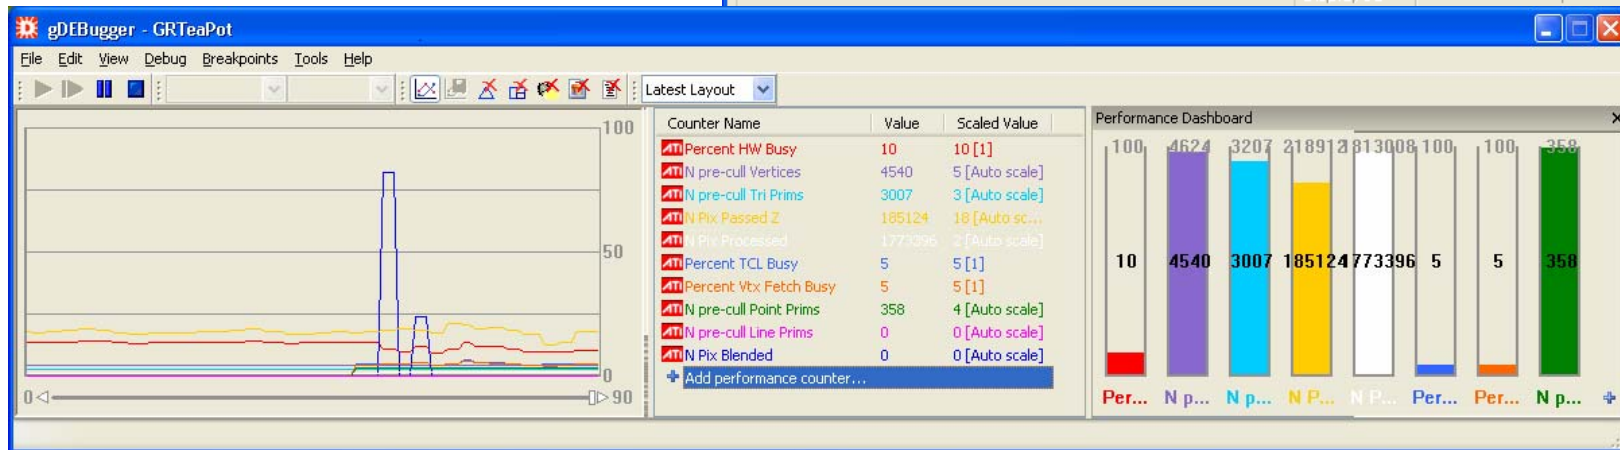
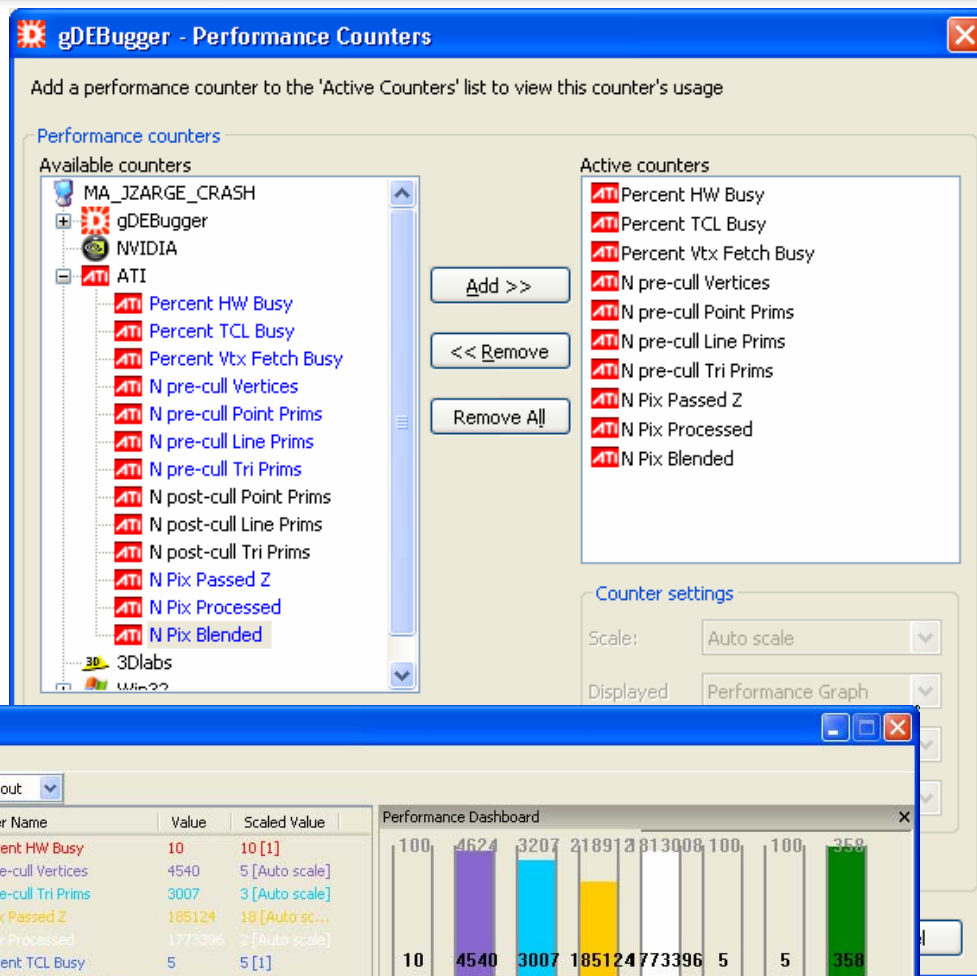

gDEBugger adds ATI performance counters

- gDEBugger & gDEBugger ES:
 - A professional OpenGL and OpenGL ES Debugger and Profiler
 - Shortens developer time required for debugging and profiling OpenGL and OpenGL ES based applications
 - Integrated with ATI performance counters to find graphic pipeline performance bottlenecks and optimize application performance
- More details:
 - www.gremedy.com
 - ATI Workstation SDK





gDEBugger Demo






The Plan

- Overview of ATI developer performance tools
- PerfDash
- ATI PIX Plugin
- gDEBugger
- **ATI content creation tools**



ATI Content Creation Tools

- RenderMonkey
 - Shader development environment
 - Supports HLSL, D3D assembler and GLSL
 - New version available soon!
- Compressorator
 - Tool for compressing textures and creating mip-map levels
- CubeMapGen
 - Tool for creating filtered cube maps without seams
 - Uses angular extent filtering
- NormalMapper
 - Automatic normal map generation tool
 - Traces rays from low resolution geometry to high resolution geometry





Wrap up

- Graphics features are a key product differentiator
- Increasing graphics efficiency can lead to
 - Playable performance on more hardware
 - More time for non-graphics functions
 - Richer stunning effects for your game
- How do you ascend to these graphic heights?
 - Incorporate the optimization techniques in to your code
 - Utilize the performance tools in your development process



Thank You!

Questions?

JZarge@ati.com

RHuddy@ati.com

devrel@ati.com

<http://www.ati.com/developer/>